

KAARA

AI USE CASE BUILD FACTORY

PRODUCT REQUIREMENTS DOCUMENT (PRD)

AI-Powered Claims Processing and Settlement — POC

with Production Requirements (POC → Production transition)

Document	PRD — POC body + Production Transition section
Use Case	AI-Powered Claims Processing and Settlement
Hypothesis ID	HYP-IN-001
Industry / Sub-industry	BFSI / Insurance
Anchor Client (Synthetic)	TrustShield Insurance
Scope — Part A (Sections 1–18)	POC release only — 10-week build
Scope — Part B (Sections 19–25)	Production requirements (POC → Production transition)
Product Owner	Kaara Vertical PM Lead
Engineering Lead	Kaara.Code Platform Engineering
Target POC Release	Week 10 (from kick-off)
Target Production v1	POC Week 10 + 14 weeks

Version / Date	v1.0 — April 2026
Status	Draft for Joint Review

Table of Contents

PART A — POC REQUIREMENTS

1. Product Overview (POC)
2. Problem Statement & Strategic Context
3. Target Users & Personas
4. POC Goals, Non-Goals, Guiding Principles
5. User Stories & Job Stories (POC)
6. POC Functional Requirements
7. UX Requirements (POC surfaces)
8. POC System Architecture & API Contracts
9. POC Data Model & Data Contracts
10. POC Integration Requirements
11. POC Security, Privacy & Compliance (baseline)
12. POC Observability, Telemetry & Logging
13. AI Model Requirements & Evaluation (POC gates)
14. POC Rollout, Experiment, and Feature-Flag Plan
15. POC Success Metrics & KPIs
16. POC Acceptance Criteria
17. Risks, Assumptions, Dependencies (POC)
18. Open Questions

PART B — PRODUCTION REQUIREMENTS (POC → Production transition)

19. Transition Overview & Phasing
20. Production Functional Extensions
21. Production Non-Functional Requirements (NFRs)
22. Production Architecture & Deployment Topology
23. Production Security, Privacy & Compliance Hardening

24. Production Observability, SRE & Run-State

25. Production KPIs and Enterprise Acceptance

Appendix

PART A — POC REQUIREMENTS

Scope: 10-week POC with TrustShield on the defined POC footprint. Everything in Part A is what gets built before the Week-10 acceptance gate.

1. Product Overview (POC)

This PRD describes the POC release of AI-Powered Claims Processing and Settlement, to be delivered over ten weeks at TrustShield Insurance. The goal of the POC is to prove — on a bounded footprint — that Kaara's AI-native solution for this use case produces measurable, executive-visible business value, while establishing the governance and data foundations required for subsequent production rollout. Part A covers what must be true at the end of Week 10. Part B covers what must additionally be true by the end of the Week 10 + 14-week production rollout.

The solution spans four capability slices corresponding to the four solution pillars in the companion Approach document, unified by a shared Compounding Memory layer. Each slice ships a minimum usable capability inside the ten weeks; full breadth is a Part B concern.

2. Problem Statement & Strategic Context

TrustShield's claims TAT averages 12 days (motor), 18 days (health); leakage estimated at 8-11%. STP rate is 14%. 480-person claims team handles rising volume at flat productivity.

Strategically, this is a compound problem. Claims pattern intelligence compounds with every settlement. This aligns precisely with Kaara's Compounding Build architecture.

3. Target Users & Personas

Persona	POC Role	Primary Need	Success by W10
TrustShield Executive Sponsor	Accountability for business outcome	Real-time outcome visibility	KPI dashboard in executive review

TrustShield Domain Lead	Governance & domain integrity	AI integrity; explainable decisions	Audit walkthrough passes
TrustShield Operational User	Day-to-day user of the AI	AI that makes my work easier, not harder	NPS $\geq +20$ from user cohort
TrustShield IT / Platform	Safe integration + operability	Clean integration; observable + reversible	Zero P1 incidents in W9–10
Internal Auditor / Compliance	Regulatory defensibility	Every decision replayable	Walkthrough pass; no P1 audit finding
External Stakeholder (e.g. customer / partner)	Experience or compliance touchpoint	Smooth, trustworthy experience	Feedback neutral-to-positive

4. POC Goals, Non-Goals, Guiding Principles

4.1 POC Goals

- Move Motor-claims TAT from 12 days to 3 days on the POC footprint by Week 10.
- Move Health-claims TAT from 18 days to 7 days on the POC footprint by Week 10.
- Move STP rate from 14% to 42% on the POC footprint by Week 10.
- Move Leakage reduction from baseline to -25% on the POC footprint by Week 10.
- Move Customer NPS on claims from baseline to +15 pts on the POC footprint by Week 10.

4.2 POC Non-Goals

- Full enterprise rollout — Part B / Production rollout.
- Adjacent use cases listed in the master tracker — out of scope for this POC.
- Full-breadth system integrations — tail systems deferred to production.
- Full autonomous action without HITL on high-value cases — envelope deliberately narrow in POC.

4.3 Guiding Principles

- AI proposes; human decides — especially at POC when trust is still being earned.
- Every POC decision is explainable and replayable; establishes the audit pattern for production.
- Data sovereignty at POC sets the norm — no customer data leaves the customer VPC.

- Default to safe — unknown or low-confidence routes to human review.
- Instrument before feature — KPIs emit from Day 1 of shadow-mode.

5. User Stories & Job Stories (POC)

5.1 Epic — Intake Intelligence

- As an operational user, I want the system to multi-channel intake (app, whatsapp, call) with instant fnol capture, damage-photo ai, and document auto-read. so I can focus on exception handling and strategy.
- As the executive sponsor, I want visibility into the business outcome this pillar drives, so I can track value capture.
- As an auditor, I want every decision from this pillar to be replayable with its rationale and model version.

5.2 Epic — Adjudication Automation

- As an operational user, I want the system to rules + ml adjudicate simple claims end-to-end; complex claims routed to adjuster with pre-work done. so I can focus on exception handling and strategy.
- As the executive sponsor, I want visibility into the business outcome this pillar drives, so I can track value capture.
- As an auditor, I want every decision from this pillar to be replayable with its rationale and model version.

5.3 Epic — Leakage Prevention

- As an operational user, I want the system to pattern detection on over-billing, inflated repair quotes, and duplicate submissions. so I can focus on exception handling and strategy.
- As the executive sponsor, I want visibility into the business outcome this pillar drives, so I can track value capture.
- As an auditor, I want every decision from this pillar to be replayable with its rationale and model version.

5.4 Epic — Settlement Memory

- As an operational user, I want the system to every settled claim, every adjudication decision, every dispute outcome enriches the model. so I can focus on exception handling and strategy.
- As the executive sponsor, I want visibility into the business outcome this pillar drives, so I can track value capture.
- As an auditor, I want every decision from this pillar to be replayable with its rationale and model version.

6. POC Functional Requirements

All requirements listed below are in-scope for the 10-week POC. Each has a priority (P0 = must-have for Week-10 acceptance, P1 = should-have, P2 = stretch) and an acceptance-test indicator. Requirements scoped only for Production are deliberately excluded from this section — see Part B, Section 20.

ID	POC Requirement	Pri	Acceptance Test
FR-P1	Ingest Claims + policy DB into Compounding Memory with normalisation to canonical domain model	P0	End-to-end ingestion test
FR-P2	Ingest Damage photos + invoices into Compounding Memory with normalisation to canonical domain model	P0	End-to-end ingestion test
FR-P3	Ingest Medical reports + hospital bills into Compounding Memory with normalisation to canonical domain model	P0	End-to-end ingestion test
FR-P4	Ingest Fraud-indicator history into Compounding Memory with normalisation to canonical domain model	P0	End-to-end ingestion test
FR-P5	Implement Intake Intelligence capability: Multi-channel intake (app, WhatsApp, call) with instant FNOL capture, damage-photo AI, and document auto-read.	P0	Functional + business-metric acceptance test
FR-P6	Implement Adjudication Automation capability: Rules + ML adjudicate simple claims end-to-end; complex claims routed to adjuster with pre-work done.	P0	Functional + business-metric acceptance test
FR-P7	Implement Leakage Prevention capability: Pattern detection on over-billing, inflated repair quotes, and duplicate submissions.	P0	Functional + business-metric acceptance test
FR-P8	Implement Settlement Memory capability: Every settled claim, every adjudication decision, every dispute outcome enriches the model.	P0	Functional + business-metric acceptance test
FR-	Return confidence score and model version with every AI	P0	API contract test

P9	decision		
FR-P10	Route high-confidence cases to automated path; low-confidence to HITL queue	P0	Queue distribution report
FR-P11	Capture every user override as labelled training signal	P0	Signal count $\geq$ threshold
FR-P12	Configurable autonomy envelope per segment and value band	P0	Config change audited
FR-P13	Operational user workbench with work-queue, rationale view, approve/edit/reject	P0	User NPS $\geq +20$
FR-P14	Auditor console to replay any POC decision with full rationale + inputs	P0	Auditor walkthrough pass
FR-P15	Executive dashboard with primary business KPI, hourly refresh, drill-down	P0	Live demo in W10
FR-P16	Model release gate: model card + bias report + rollback plan per release	P0	Gate artefacts stored
FR-P17	Admin console for roles, autonomy envelopes, integrations, release gates	P0	Smoke test in W10

7. UX Requirements (POC surfaces)

7.1 POC Surfaces

- Operational Workbench (web) — daily driver for the operational user role.
- Auditor Console (web) — replay and verify any POC decision.
- Executive Dashboard (web; mobile-responsive) — primary business KPI with drill-down.
- Admin Console (web) — roles, thresholds, integration management, release gate.

7.2 Design Tenets

- Decision-first, not feature-first.
- Rationale alongside every answer.

- One-click override with required reason code.
- WCAG 2.1 AA; keyboard navigation; high-contrast mode.

8. POC System Architecture & API Contracts

8.1 POC Logical Architecture

The POC architecture uses the six-layer reference design (source systems → ingestion → Compounding Memory → AI services → workflow/HITL → governance/observability) sized for POC load. Production hardening, HA clusters, and multi-region topology are Part B (Section 22).

8.2 POC Core APIs

API	Method + Path	Purpose (POC)
Event Ingest	POST /v1/events	Push a new domain event into the platform
Decision Propose	POST /v1/decisions/propose	Request AI-proposed decision for a given input
Decision Execute	POST /v1/decisions/execute	Execute an approved decision via the right action surface
Override	POST /v1/decisions/{id}/override	Capture a user override with reason code
Decision Replay	GET /v1/decisions/{id}	Return full rationale + inputs + model card
Autonomy Config	GET/PUT /v1/config/autonomy	Read or update the autonomy envelope

9. POC Data Model & Data Contracts

Canonical entities in POC scope include: source-domain entities from the customer systems (normalised to canonical form), decision entities (inputs, AI output, rationale, confidence, model version), override entities, model-artefact entries in the registry, and audit-event

entries. Each ingestion source has a contracted schema, freshness SLA, and owner; PHI/PII fields are tagged at schema level and auto-tokenised.

10. POC Integration Requirements

System	POC Status	Mechanism
Guidewire ClaimCenter	In scope (POC)	REST / file-drop / event stream as applicable
Majesco Policy	In scope (POC)	REST / file-drop / event stream as applicable
Network-hospital APIs (health)	In scope (POC)	REST / file-drop / event stream as applicable
Workshop-estimate platforms (motor)	In scope (POC)	REST / file-drop / event stream as applicable
UPI + NEFT for payouts	In scope (POC)	REST / file-drop / event stream as applicable

Integrations deferred to Part B / Production: tail systems, full ABDM/HCX / equivalent regulator flows, SIEM forwarding, and deep bidirectional write-back where not required for POC acceptance.

11. POC Security, Privacy & Compliance (baseline)

POC-level security is production-grade for the essentials: encryption at rest and in transit, authentication and access control, tokenisation of identifiers, and immutable audit logs. Performance, scalability, and availability targets are deferred to Part B — see Section 21. This section captures only the security-and-privacy baseline required for POC.

11.1 Security Baseline

- All services in private subnets; egress allow-list only.
- mTLS between services; service identities via SPIFFE.
- Secrets in HashiCorp Vault; 90-day rotation minimum.
- SSO via Okta with MFA on admin roles.
- All authN/authZ and data access logged to tamper-evident store.

11.2 Privacy Baseline

- DPDP 2023 compliant — consent capture, purpose limitation, breach-notification SOP.
- Identifiers tokenised at ingestion; reverse-lookup audited.
- Privacy-impact assessment completed by Week 2.
- DPO engaged from Week 1; DPO walkthrough scheduled for Week 9.

11.3 POC Compliance Evidence (light-touch)

Full certification / attestation scope (ISO 27001, SOC 2, sector-specific certifications) is Part B. POC-level evidence includes: access-control logs, audit trail for every decision, DPO sign-off, and a written POC risk assessment.

12. POC Observability, Telemetry & Logging

12.1 Technical Telemetry (POC)

- Golden signals per service: latency, traffic, errors, saturation.
- End-to-end trace ID per decision; traces visible in Tempo.
- POC uptime dashboard shared daily with customer CTO office.

12.2 Business Telemetry (POC)

- Decisions processed per day, by state (auto / HITL / exception).
- Core business metric trending vs baseline.

- User workbench usage telemetry.
- Compounding Memory growth — learning signals captured per week.

12.3 Governance Telemetry (POC)

- Weekly bias-probe run against curated probe-set.
- Override-rate monitor; spikes flagged for review.
- Model drift signal — input and output distribution, weekly.

13. AI Model Requirements & Evaluation (POC gates)

POC model gates are the 'shadow' and 'production-band' thresholds required for the POC environment. Full production-promotion gates are in Part B Section 25.

Model (per pillar)	Class	POC Shadow Gate	POC Production-Band Gate
Intake Intelligence Model	Classical ML + LLM as applicable	≥88% expert-agreement on blinded sample	≥92% expert-agreement + no safety incident for 2 weeks
Adjudication Automation Model	Classical ML + LLM as applicable	≥88% expert-agreement on blinded sample	≥92% expert-agreement + no safety incident for 2 weeks
Leakage Prevention Model	Classical ML + LLM as applicable	≥88% expert-agreement on blinded sample	≥92% expert-agreement + no safety incident for 2 weeks
Settlement Memory Model	Classical ML + LLM as applicable	≥88% expert-agreement on blinded sample	≥92% expert-agreement + no safety incident for 2 weeks

13.1 Evaluation Methodology (POC)

- Blinded weekly sample audit — 2 senior domain experts + 1 external auditor.
- Bias probes — per relevant protected attribute and business segment, re-run weekly.
- Counterfactual eval on overrides — check overridden state against held-out labels.
- A/B control on new model versions — new version gets 10% traffic until gate is met.
- Safety gate — high-value and flagged-segment cases always route HITL during POC regardless of confidence.

14. POC Rollout, Experiment, and Feature-Flag Plan

Week	Flags flipped / scope
W5	Models in shadow (historical); sub-set of integrations live
W6	Remaining pillars shadow; appeal / exception flows drafted
W7	All capabilities shadow on live traffic; zero-action
W8	User console live; users approve AI proposals with override capture; daily shadow metrics
W9	Production-band cut-over: bounded volume slice; narrow autonomy envelope
W10	Production-band expanded where safe; demo + acceptance gate

15. POC Success Metrics & KPIs

KPI	Baseline	POC Target (W10)
Motor-claims TAT	12 days	3 days
Health-claims TAT	18 days	7 days
STP rate	14%	42%
Leakage reduction	baseline	-25%

Customer NPS on claims	baseline	+15 pts
------------------------	----------	---------

16. POC Acceptance Criteria

POC is accepted when all 'must-have' criteria below are met as of Week 10, jointly verified by Kaara and customer. Stretch criteria are non-binding but trigger accelerated production timeline.

ID	Criterion	Must / Stretch	Evidence
AC-1	Motor-claims TAT at 3 days on production-band volume	Must	W9-10 dashboard export
AC-2	Health-claims TAT at 7 days on production-band volume	Must	W9-10 dashboard export
AC-3	STP rate at 42% on production-band volume	Must	W9-10 dashboard export
AC-4	Leakage reduction at –25% on production-band volume	Must	W9-10 dashboard export
AC-5	Customer NPS on claims at +15 pts on production-band volume	Must	W9-10 dashboard export
AC-6	Every POC decision is auditable and replayable	Must	Auditor walkthrough
AC-7	Executive dashboard live, updated hourly, drill-down working	Must	Live demo + telemetry
AC-8	Operational user NPS $\geq +20$	Must	Survey
AC-9	DPO walkthrough completed with no P1 findings	Must	DPO attestation
AC-10	Production-readiness report signed off by CTO	Must	Signed report

AC-11	Bias probe parity within agreed threshold across tested cohorts	Must	Weekly probe report
-------	---	------	---------------------

17. Risks, Assumptions, Dependencies (POC)

17.1 Key POC Risks

- Image-based damage assessment accuracy (med; high) — HITL on low-confidence; continuous-training with adjuster-labelled corrections.
- IRDAI regulatory scrutiny on STP (med; high) — Explainable decisions; TAT compliance; complaint-ratio monitoring.
- Provider (hospital/workshop) pushback (med; med) — Transparent rationale; dispute-resolution SLAs; continued provider-engagement.

17.2 POC Assumptions

- TrustShield provides signed data-use and access agreements by Week 1.
- Executive sponsor and domain lead available for the 10-week window.
- Sufficient volume in the POC footprint for statistically-meaningful signal.

17.3 POC Dependencies

- Access to Guidewire ClaimCenter confirmed by customer IT by Week 2.
- Access to Majesco Policy confirmed by customer IT by Week 2.
- Access to Network-hospital APIs (health) confirmed by customer IT by Week 2.

18. Open Questions

1. Is the POC autonomy envelope acceptable to the customer Operating Council, or does it need tightening for specific segments?
2. What is the acceptable external-stakeholder disclosure language (e.g. 'AI-assisted') for this use case?
3. Should the POC footprint include a challenging edge segment for realistic bias-probe testing, even if it raises complexity?
4. How should success-fee be pro-rated if only some must-have criteria are met?
5. Does TrustShield consent to this POC being used as a public case study post-acceptance?

PART B — PRODUCTION REQUIREMENTS

(POC → Production transition)

Scope: everything additionally required for production rollout across the full enterprise footprint within 14 weeks of POC acceptance. Developers and end-users use Part B to plan what changes between POC and full-scale deployment.

19. Transition Overview & Phasing

Part B describes the delta. Every requirement in this section describes how the POC build must harden, scale, or extend when moving from the POC footprint to the full customer footprint. The rollout is phased over 14 weeks so that scale is added only after the POC-proven slice is hardened.

Phase	Weeks	Focus	Exit Gate
Hardening	1–3	Performance, DR, secrets, access review, VAPT	Infra audit sign-off; DR drill successful
Scale wave 1	4–6	Expand from POC slice to ~50% of footprint	Core metric trending on plan
Scale wave 2	7–10	Extend to full footprint	All entities live; full payer/supplier coverage
Stabilise	11–12	Noise reduction, feedback loops, user upskilling	Steady-state production quality
Handover	13–14	Transition to run-state; model-refresh calendar; commercials	Formal handover; Y1 run-state begins

20. Production Functional Extensions

These are requirements added on top of the POC functional set when moving to production. Each extension builds on the POC foundation laid in Part A.

ID	Production Extension	Pri
FR-R1	Extend ingestion + integrations to full enterprise footprint (all entities / business units)	P0
FR-R2	Full payer / supplier / partner API coverage; RPA fallback for tail systems	P0
FR-R3	Full regulator-facing flows (sector-specific) — ABDM HCX, SEBI, RBI as applicable	P0
FR-R4	Deep integration surface — e.g., physician nudge in EHR, trader copilot in OMS, etc.	P1
FR-R5	Bulk re-scoring of open cases after model refresh	P1
FR-R6	Ecosystem / partner dashboards (e.g., payer-twin, supplier-twin, advisor-twin)	P1
FR-R7	Adjacent-use-case hooks enabling Year-2 adjacencies	P2
FR-R8	What-if simulation surface for business leadership	P2

21. Production Non-Functional Requirements (NFRs)

NFRs apply to the production environment. POC NFRs are intentionally not required; the POC is sized for proving business value on a bounded footprint, not for meeting enterprise-scale non-functional targets.

Category	Production NFR	Production Target
Performance	AI decision inference latency (p95)	≤3 seconds
Performance	End-to-end decision cycle (autonomous band)	Use-case specific; target

		≤3 hours
Scalability	Concurrent decisions processed	Use-case specific; sustained + burst targets tracked
Availability	System uptime — business hours / overall	99.5% / 99.0%
Availability	DR — Recovery Point / Recovery Time Objective	≤15 min / ≤4 hours
Availability	Region topology	Primary region + warm DR region; 3-AZ primary
Security	Encryption at rest / transit	AES-256 / TLS 1.3
Security	Authentication	SSO + MFA; admin MFA mandatory
Security	Authorisation	RBAC + ABAC; role segregation across users / reviewers / auditors
Privacy	DPDP data-principal rights console	Access, rectification, erasure, portability — functional
Auditability	Retention	7 years, immutable, tamper-evident
Maintainability	Deployment frequency	Weekly; zero-downtime rolling deploys
Maintainability	Portability	Cloud-agnostic architecture; artefacts reproducible on Azure/GCP

22. Production Architecture & Deployment Topology

22.1 Topology

- Region — primary + DR; 3-AZ in primary.
- Compute — EKS with HPA; Fargate for inference burst; managed stateful stores (RDS Multi-AZ, MSK 3-broker, OpenSearch 3-node, Neo4j Enterprise).
- Network — zero-trust service mesh (SPIFFE/mTLS); private VPC endpoints for S3, Bedrock, KMS.
- IaC — Terraform + Terragrunt; environment parity (dev/stg/prod); policy-as-code via OPA.
- Release — ArgoCD GitOps; canary deploys; automated rollback on SLO-burn-rate breach.
- Model registry — MLflow; promotion requires signed model card + bias report + rollback plan.

22.2 Scale Testing

- 3x peak-volume synthetic load tests before each scale-wave release.
- Randomised chaos experiments against stateful dependencies.
- End-to-end trace-level performance profile at every scale wave.

23. Production Security, Privacy & Compliance Hardening

23.1 Security Hardening

- Quarterly third-party VAPT; annual certification.
- Daily SCA; weekly container scan; monthly SAST; continuous dependency audit.
- SIEM forwarding over TLS; 24x7 SOC cover (Kaara-managed).
- Tenant isolation — dedicated AWS account per customer; no shared compute or storage.
- Quarterly principle-of-least-privilege access reviews.

23.2 Privacy Hardening

- DPDP data-principal-rights console live; SLAs on access / rectification / erasure documented.
- DPO on-call rota; breach-notification SOP exercised quarterly.
- Consent artefact lifecycle — captured, retained, and retractable.

23.3 Compliance Certifications (Y1 roadmap)

- ISO 27001 — certification within 9 months of production go-live.
- SOC 2 Type II — attestation within 12 months of production go-live.
- Sector-specific certifications (e.g., NABH, HIPAA-alignment for healthcare; PCI-DSS or RBI-alignment for BFSI) as applicable.

24. Production Observability, SRE & Run-State

24.1 SLOs & Error Budgets

- SLOs for latency, availability, and business correctness.
- Burn-rate alerts on 1h / 6h / 24h windows — not raw threshold alerts.
- Monthly error-budget reviews; feature velocity slows when budget is burnt.

24.2 Run-State Team

Function	Role	FTE
Platform SRE	24x7 rota, alerting, incident mgmt	4.0 (distributed)
Model Operations	Drift monitoring, weekly retraining	1.5
Domain Operations	Shadow audit, exception triage, feedback	2.0 + embedded
Account & governance	Monthly governance council, quarterly review	0.5
Compliance & audit	Regulatory obligations, evidence pack	0.5

24.3 Production SLAs (customer-facing)

- Decision-to-outcome SLA: use-case specific (see Section 21).
- Platform uptime: 99.5% business hours; 99.0% overall.
- P1 incident: 15 min ack, 2h resolution target, RCA ≤72h.
- Model drift: weekly review; retraining ≤10 business days of threshold breach.
- Security incident: DPO ≤4h, affected-parties notification per DPDP norms.

25. Production KPIs and Enterprise Acceptance

25.1 Year-1 KPIs (Production)

KPI	POC Target	Production Y1 Target
Motor-claims TAT	3 days	≤24 hrs
Health-claims TAT	7 days	≤3 days
STP rate	42%	≥65%
Leakage reduction	–25%	–50%
Customer NPS on claims	+15 pts	+30 pts

25.2 Enterprise Acceptance Criteria (Y1)

- Full enterprise footprint live on the solution stack.
- Core business metric at the Production-Y1 target.
- Autonomous decisioning share at Production-Y1 target within agreed envelope.
- Executive sponsor publishes quarterly outcome dashboard at the executive committee or board.
- Zero P1 incidents from Part B go-live through Y1 anniversary.
- ISO 27001 certification roadmap on-track; SOC 2 Type II scheduled.

Appendix

Linked Documents

- POC Approach & Productionisation Plan (companion document).
- Kaara.Code Platform Reference.
- TrustShield Insurance synthetic datasets — claims.csv, policies.csv.
- DPIA template — Kaara vertical specific.
- AI_UseCases_Master_Tracker_Kaara.xlsx — Hypothesis HYP-IN-001.